

10 Funding Applications

A Phase 1, First-in-Human, PDAC Clinical Trial Protocol of a
LLM-Directed Robotic Whipple with Daraxonrasib (RMC-6236)

Draft 1.0

Repository v4.4.0

20 figures

10 recipients

01 NIH Pioneer Award

02 ARPA-H

03 NSF TIP X-Labs

04 DOE Genesis Mission

05 NIH SEED SBIR

06 FNIH AMP

07 HHMI Investigator

08 NCI CTEP

09 Convergent FRO

10 UC San Diego Moores

CEO Kevin Kawchak , ChemicalQDevice | kevink@chemicalqdevice.com
[10.5281/zenodo.21787424](https://doi.org/10.5281/zenodo.21787424) | 0009-0007-5457-8667
San Diego, CA | August 4, 2026

Independent research paper and practical adoption guide. It is not medical or regulatory advice and is not endorsed by the FDA, NIH, HHS, an IRB, ICH, or any sponsor. All figures derive from the author's repository sources and are illustrative unless tied to a cited reference.

Disclaimer: This work is independent and is not endorsed or sponsored by any trial sponsor, CRO, site, IRB, regulator, or medical society; and was adapted using Claude Code Opus 5.

The paper (v1.0) is at <https://doi.org/10.5281/zenodo.21787424> and the repository (v4.4.0) is at <https://github.com/kevinkawchak/physical-ai-oncology-trials>. The ten application file sets summarised here are in [funding/pdac-funding-applications](#).

Keywords Independent scientist; Federal research portfolio; Pancreatic ductal adenocarcinoma; Robotic pancreaticoduodenectomy; Daraxonrasib; Physical AI; On-premises LLM; Phase 1 trial; Funding mechanisms

Table of Contents

Abstract	3
How to Read This Paper	3
Figure and Table Inventory	3
1 The \$200 Billion Realignment and the Individual Scientist	5
1.1 What the Report Actually Says	5
1.2 Where the Enterprise Stalls	6
1.3 Why an Independent Scientist Is the Test Case	6
2 The Ten Applications	7
2.1 Selection Rule	7
2.2 What Differs, and What Does Not	8
2.3 Redundancy Across the Set	8
3 Set A: The Surgical Perspective	10
3.1 What Leads in Set A	10
3.2 Who Is Authorised to Do What	10
3.3 Applications 01 to 05	11
3.4 What a Reviewer Actually Asks	12
4 Set B: The Medical Oncology Perspective	13
4.1 What Leads in Set B	13
4.2 Applications 06 to 10	13
4.3 What Happens If Any One Lands	13
4.4 The Two Sets Describe One Operation	14

5	Trial Evidence	15
5.1	The Chronology	15
5.2	What the Numbers Are, and What They Are Not	16
5.3	Cost and Schedule	17
5.4	Provenance	17
6	Physical AI Governance	18
6.1	Three Boundaries	18
6.2	The Operative Day	18
6.3	How a Failure Propagates	18
6.4	What Is Released	18
6.5	Who Produced Which Artifact	21
7	UC San Diego Moores Cancer Center	22
7.1	The Positioning Constraint	22
7.2	Three Corrections the Site Research Forced	22
7.3	Two Routes, Sequenced Differently	23
7.4	What a Productive Early Outcome Looks Like	23
8	Budget and Leverage	23
8.1	One Budget, Ten Shapes	23
8.2	Leverage	25
8.3	The Binding Constraint Is Not Money	25
9	Build Method	26
9.1	Two Schedules, Thirteen Sub-Prompts	26
9.2	Why the Diagram Split Is Uneven	26
9.3	One Operator Against a Program Office	27
9.4	What Is Reproducible, and What Is Not	27
10	Risks and Limits	28
10.1	What Is Not Claimed	28
10.2	What Would Falsify the Case	29
10.3	Risks the Programme Cannot Retire Alone	29
10.4	The Limit of the Whole Exercise	29
	Back Matter	30
	Abbreviations	30
	Data and Code Availability	30
	Author Contributions and Conflicts	30
	Citation	30
	References	31

Abstract

In July 2026 the White House asked the federal research portfolio, approximately \$200 billion annually, to prioritise the individual scientist over legacy institutions [1]. This paper reports what happens when one unaffiliated scientist takes that literally and writes ten funding applications for the same Phase 1 pancreatic cancer trial, each addressing the mechanism its recipient runs.

The trial is constant across all ten: an open-label, single-arm, 3+3 escalation of perioperative daraxonrasib (RMC-6236) in up to 18 participants with resectable or borderline-resectable KRAS G12-mutated pancreatic ductal adenocarcinoma, performed as a staged eight-arm robotic pancreaticoduodenectomy with an on-premises language model confined to advisory output [2, 3]. Five lead with the operation and five with the drug; both sets describe the same hybrid procedure. UC San Diego Moores Cancer Center is named in all ten as the intended partner of choice, at the feasibility stage only, and described nowhere as a partner, sponsor, site, or endorser.

The evidence is stated with its limits attached. An August 2025 quantitative systems pharmacology simulation returned median overall survival of 12.8 months against 5.4 for chemotherapy [4]; in May 2026 RA-Solute 302 reported 13.2 against 6.6 months in the RAS G12 metastatic population [5]. That proximity is chronology, not validation, and neither result speaks to resectable setting, which is what trial would measure.

The applications carry no digital object identifiers. This paper carries one.

How to Read This Paper

This is a summary of ten application file sets, not an application. Every claim in it can be checked against a named file in a public repository [6].

It is not medical or regulatory advice, and not a claim that any institution, sponsor, or drug developer has agreed to anything. No patient has been treated. The trial does not exist until a site, an institutional review board, an active investigational new drug application, monitoring, insurance, and a sponsor of record are all in place.

Three reading paths. A **funder** should read §2 and §8. A **clinician** should read §6 and §5. A **metascience reader** should read §1 and §9.

Figure and Table Inventory

Twenty figures across five machine-readable diagram platforms. The split follows what each figure has to answer rather than an equal quota: an even division would have forced at least four figures into a vocabulary that cannot state their subject.

Fig	Type	Sec	The question it answers
1	mermaid	§1	How does one policy sentence become ten addressed applications?
2	plantuml	§1	Where does unaffiliated proposal stall, and what removes each stall?
3	d2	§2	How do the ten compare on mechanism, term, ask, and lead?
4	graphviz	§2	Which award unblocks which activity, and what is redundant?
5	plantuml	§3	Who is authorised to take which action?
6	mermaid	§3	What are the reviewer's gates, and which section answers each?
7	plantuml	§4	What runs in parallel once any one award lands?
8	mermaid	§5	How long did each activity run, and what overlapped?
9	d2	§5	Which evidence tier contains which, and what may not be cited?
10	graphviz	§5	Which prior work fed which, and where do the two lines meet?
11	diagrams	§6	What runs inside the trust boundary, and what may cross it?
12	mermaid	§6	Who says what to whom across one operative day?
13	plantuml	§6	Under what guard does the advisory system change state?
14	graphviz	§6	By what path could an unintended motion reach the patient?
15	d2	§6	What joins to what in the released dataset?
16	diagrams	§6	Which model produced which artifact, and who reviewed it?
17	d2	§8	Where does each dollar land, and what is contributed vs. paid?
18	mermaid	§8	How long is each review clock against the site milestone?
19	diagrams	§9	How does one operator's stack compare with a program office?
20	graphviz	§10	Which field can falsify the programme's central claim?

Table 1. The twenty figures, their platform, their section, and the question each exists to answer. Five are mermaid-type, four each plantuml-type, d2-type and graphviz-type, and three diagrams-python-type. Section 7 carries no figure: application 10 already holds its intake sequence and authority grid.

Table	Sec	Contents	Source directory
1 to 2	Front	Figure and table inventories	The five diagram stage directories
3 to 4	§1	Seven findings & correctives; goals answered	funding/science-golden-age/
5	§2	Ten recipients and the sentence qualifying each	applications/
6	§3	Set A figure inventory, including type not used	applications/app-01 to app-05
7	§4	Consortium measurements vs. resection timing	applications/app-06
8 to 9	§5	Four results, stated limitations; cost & schedule	funding/supplementary/source-files/
10	§6	The three frontier-model roles	funding/tripartisan-llm-support.md
11 to 12	§7	3 positioning corrections; 2 site routes	funding/potential-partners/
13	§8	One budget in ten shapes	funding/RFA-RM-27-001-v2/
14 to 15	§9	Thirteen sub-prompts; the diagram split	sub-prompts/
16 to 17	§10	What isn't claimed; risks programme can't retire	Application back matter
18	Back	Abbreviations at the body measure	Published Phase 1 protocol, IND

Table 2. The paper's eighteen tables, the section each appears in, and the repository directory its values are drawn from. No value in any table is computed here for the first time. Tables 1 to 2 sit in the front matter, Table 18 in the back matter, and neither position carries a section number.




Figure 1. One sentence of federal policy, three mechanism families, and the ten recipients each family reaches. The split is by mechanism, not by agency, which is why one recipient is not a federal funder.

1 The \$200 Billion Realignment and the Individual Scientist

1.1 What the Report Actually Says

The July 2026 report to the President, *Science: A New Golden Age*, states four overarching goals [1]: prioritise the individual scientist over legacy institutions; change how research dollars are allocated, distributed and assessed; set clear scientific goals and build capacity to translate them; and prepare the enterprise for the AI revolution.

The first is stated without hedging: “If we are serious about expanding what our scientists can do, we must invest directly in American researchers and the bold ideas that drive them. Too much of our research enterprise has come to serve itself rather than the scientists within it.” The Summary turns that into an instruction with a named model attached: “bet on people, not just projects,” by “scaling long-horizon grants for the best and brightest modeled on National Institutes of Health (NIH) Director’s Pioneer Award.”

Chapter II attaches a number. Federal agencies distribute approximately **\$200 billion** in annual research and development funding, and the finding about that portfolio is the sentence this paper is written against: there is “no systematic framework for identifying where those dollars could catalyze the greatest scientific returns, with deference instead given to the same incumbents that consume the funding.”

The corrective has three parts: realign toward risk taking, strip away the administrative burden that consumes close to half of an investigator’s working hours, and embed continuous evidence-based improvement. What the report does not do is say who should receive the realigned money. This paper is one answer, submitted ten alternate times.




Figure 2. The five states a proposal from an unaffiliated scientist passes through, and the three at which the report says the enterprise currently stalls. Each stall names the mechanism that removes it.

1.2 Where the Enterprise Stalls

The report's diagnosis is quantitative, and every number below is the report's own or is drawn from the primary source its end notes name.

Finding	Value	Mechanism the report names against it
Research time lost to administration	near one half	Burden reduction; applications of a few pages
New federal requirements on grants, 1991 to January 2025	at least 270	Requirement rescission; metascience units
Negotiated indirect cost rates	50-60 percent	Indirect-cost reform; private funders accept 10 to 15
Grant lead time	≤ 20 months	Fast grants decided in under one month
NIH principal investigator average age, 1980 to 2008	39 rising to 51	Early-career independence; portable fellowships
Fall in new drugs per billion dollars since 1950	roughly eightyfold	Portfolio construction across exploration and exploitation [7]
Publication-rate advantage of person-based funding	nearly 2x	HHMI-style support, roughly \$10M over seven years [8]

Table 3. The report's diagnosis in seven numbers, each paired with the mechanism it names as the corrective. The last row is the only one that reports an outcome rather than a friction.

An unaffiliated scientist meets a subset of these frictions and one the table does not contain: the eligibility screen. Most federal mechanisms require a host institution before a proposal can be reviewed at all, so the merit question is never reached. Figure 2 states the resulting state machine and names, for each stall, the mechanism the report offers against it.

1.3 Why an Independent Scientist Is the Test Case

Between June 2025 and July 2026 one person produced fourteen deposited works: two clinical protocols, an investigational new drug application, two funding applications, five legislative drafts, and four guidance works [9], none carrying indirect cost. Three trial simulations were completed at an estimated \$36,330 each against an industry benchmark running from \$120,000 for a single empirical virtual trial to more than \$2 million for a ten-arm quantitative systems pharmacology trial [10].

The claim is a ratio and nothing more. Not that the work is better than an institution's, and not that institutions are unnecessary: only that the part of the pipeline which used to require a program office no longer does, and that the part which still requires an institution, patients and a theatre and a review board, is exactly the part these ten applications are asking to reach.

d2-type / scored grid

Recipient	Mechanism	Term	Direct ask	Lead
01 NIH Pioneer Award	person-based	5 years	\$3.5M	surgical
02 ARPA-H	milestone	36 months	\$2.1M	surgical
03 NSF TIP X-Labs	organization	5 years	\$3.5M	surgical
04 DOE Genesis Mission	mission	5 years	\$3.5M	surgical
05 NIH SEED SBIR	small business	9 + 24 mo	\$1.6M	surgical
06 FNIH AMP	consortium	5 years	\$3.5M	medical
07 HHMI Investigator	person-based	7 years	\$0.7M/yr	medical
08 NCI CTEP	concept review	5 years	\$3.5M	medical
09 Convergent FRO	time-bound	5 years	\$3.5M	medical
10 UC San Diego Moores	institutional	one meeting	none	medical

The science column is omitted because it would be constant: the same population, the same 3+3 escalation, the same co-primary endpoints, and the same advisory boundary appear in all ten.

Figure 3. Ten applications on four attributes. The science column is constant by construction: the same trial is described ten times, and only the mechanism, the term, and the ask change.

Transmittal-letter goal	What the ten applications supply against it
Prioritise the individual scientist over legacy institutions	One named applicant, no host institution, fourteen deposited works before any award
Change how research dollars are allocated and assessed	Ten mechanism families addressed on the same science, so the mechanism is the only variable
Set clear goals and build translational capacity	A published protocol and IND, and a governance package designed to be site-portable
Prepare the enterprise for the AI revolution	An advisory system whose boundary is measured in milliseconds and released as an open schema

Table 4. The report's four overarching goals, and the specific thing the application set supplies against each. The second row is what makes the set a comparison rather than ten separate requests.

2 The Ten Applications

2.1 Selection Rule

A recipient qualifies only if *Science: A New Golden Age* names its programme or the mechanism that programme runs. Nothing was included on plausibility. The rule is what makes the set an argument rather than a mailing list: every application opens by quoting the recipient's own mechanism back to it, in the administration's wording, and then answers it.

Three families emerged: **person-based**, where the report names the Pioneer Award by title and cites HHMI-style support as the evidence; **organization-type**, where it names X-Labs, focused research organizations, the ARPA program-manager model and SBIR; and **partnership and cost share**, where it holds up FNIH and the Accelerating Medicines Partnership as its model, names the Genesis Mission's Robotics line, and asks agencies to prioritise non-federal cost share [11].

2.2 What Differs, and What Does Not

No	Recipient	Mechanism	The anchor sentence in the report
01	NIH Common Fund, Director's Pioneer Award	Person-based	"scaling long-horizon grants for the best and brightest modeled on . . . Director's Pioneer Award"
02	ARPA-H mission office	Organization-type	The ARPA program-manager model, with ARPA-H and ARPA-E as the live examples
03	NSF TIP Directorate, X-Labs	Organization-type	"the first federal program designed to fund independent research organizations outside traditional academia"
04	DOE Office of Science, Genesis Mission	Partnership	Robotics mission: "to initiate the era of physical AI-driven scientific discovery"
05	NIH SEED, SBIR and STTR	Organization-type	"Programs like SBIR open doors for technician-founded ventures"
06	Foundation for the NIH, AMP	Partnership	FNIH and the Accelerating Medicines Partnership held up as the model
07	HHMI Investigator Program	Person-based	Person-based funding of roughly \$10 million over seven years
08	NCI Cancer Therapy Evaluation Program	Partnership	Prioritised foundational biological sciences; the report's own cancer framing
09	Convergent Research, FRO programme	Organization-type	Focused research organizations as time-bound nonprofit research startups
10	UCSD Moores Cancer	Partnership	Regional innovation clusters; non-federal cost share

Table 5. The ten recipients with the sentence in the report that qualifies each. No recipient appears without an anchor, which is what makes the set an argument inside the funder's own frame rather than a mailing list.

What is constant is the trial. All ten describe the same population, the same open-label single-arm 3+3 escalation to 18 treated participants, the same co-primary endpoints, and the same advisory boundary at 3 milliseconds per arm and 500 milliseconds system-wide [2, 3].

That constancy is the point. A reviewer comparing two applications from this set is comparing two funding mechanisms applied to one piece of science with the science held fixed. It is the closest thing an individual applicant can build to a controlled comparison of mechanisms, and it is possible only because the protocol was published before the applications were written.

2.3 Redundancy Across the Set

Sending ten applications is not ten times one application. Because they fund overlapping activity sets, most of the programme survives any single refusal. Figure 6 makes that concrete and identifies the two activities that do not.

Six of the eight activities have two or more upstream awards: the interlock rig from applications 02, 03 and 05; the logging schema from 03, 04, 05 and 09; cohort accrual from 01, 02, 06, 07 and 08.

The two exceptions are regulatory. IND activation is reachable only through applications 01 and 02, and IRB approval only through application 10, which asks for no money at all. That asymmetry is the programme's real structure: the scarce resource is institutional standing, not funding, and the one application that could supply it is the one with nothing to offer in return.




Figure 4. Ten awards against eight trial activities. Six activities have two or more upstream awards and survive any single refusal; the two that do not are drawn heavier, and both are regulatory.




Figure 5. Seven actors and the eleven actions each may take. Every action has at least one human association; the advisory model has exactly one, and the struck links are the four it is denied by construction.

3 Set A: The Surgical Perspective

3.1 What Leads in Set A

In the five surgical applications the operation leads. Daraxonrasib enters as the perioperative arm the operation is timed around rather than as the headline, and the endpoint table is led by conversion, anastomotic integrity, operative time, blood loss, R0 margin, and 90-day morbidity [12].

That ordering is not cosmetic. A mechanism office asks whether the hard part is buildable; a clinical programme asks whether it is survivable. Set A answers the first, and the hard part it names is the operation with a language model in the room.

3.2 Who Is Authorised to Do What

The most common objection to a language model in an operating theatre is that the boundary is asserted rather than specified. Applications 02 and 10 each carry a three-actor slice. Figure 5 gives the whole set, including the two actors no application figure contains: research pharmacist and institutional review board.

Three fills separate three authority classes: Corporate Blue for the actions only the operating surgeon takes, pale blue for shared clinical and oversight actions, gray for the three the model may take, all of which produce text or read data and none of which moves anything. A reader who ignores the labels still reads the structure correctly.

3.3 Applications 01 to 05

01, NIH Director’s Pioneer Award. The unit of risk is the individual, not the aim. The bet does not survive decomposition into three project aims, because the drug schedule, the operative plan, and the advisory boundary are only interesting together. Update: 7/6/26 emailed Pioneer Award RFA-RM-27-001 v1 to NIH Director. Application v2 is complete, public, and ready for 9/9/26 submission deadline.

02, ARPA-H. A gated programme rather than a project narrative, because the report endorses the ARPA model on the ground that a program manager can end a bet. Three gates, each with a numeric kill criterion checkable from programme data rather than from a progress report. Update: 8/2/26 applied to ARPA-H.

03, NSF TIP X-Labs. Argues organization type rather than money. Five activities must be held together and no existing performer type covers all five, which is the gap the report says X-Labs was created to fill.

04, DOE Genesis Mission. The FY 2028 annex defines the Robotics mission as “general-purpose autonomous systems capable of dexterous manipulation, mobility, and reliable operation in real-world environments, to initiate the era of physical AI-driven scientific discovery” [11]. A human operating room is the hardest available instance, and three artifacts are offered to mission rather than to clinic [13].

05, NIH SEED SBIR. The only one where the commercial product carries the argument. The deliverable is not the robot and not the drug but the governance package that makes an on-premises model usable.




Figure 6. The five questions a reviewer asks in order, and the section of the application that answers each. A no at gate 1 or gate 2 ends the review; a no at gates 3 to 5 changes the ask rather than ending it.

3.4 What a Reviewer Actually Asks

The asymmetry between the first two gates and the last three is why every attachment settles credibility and plausibility in its first two sections, before any ask: a reviewer unconvinced there never reaches budget.

No	Figures	Types used	Type deliberately not used, and why
01	3	mermaid, d2, graphviz	No plantuml or diagrams: no authority question, no topology question
02	4	mermaid, plantuml, d2, graphviz	No diagrams: nothing is deployed across a boundary in this argument
03	3	d2, diagrams, mermaid	No plantuml or graphviz: the argument is organization type, not authority or dependency
04	4	diagrams, mermaid, graphviz, d2	No plantuml: the guards are stated in prose and drawn in the summary paper instead
05	3	mermaid, graphviz, d2	No plantuml or diagrams: a phased commercial argument needs neither

Table 6. Figure inventory for Set A, including the platform each application deliberately did not use. Recording the omission is what keeps the type choice a decision rather than a default.

4 Set B: The Medical Oncology Perspective

4.1 What Leads in Set B

In the five medical oncology applications the drug and the patient selection lead. Daraxonrasib (RMC-6236) is a RAS(ON) multi-selective inhibitor given perioperatively in KRAS G12-mutated resectable or borderline-resectable disease, and the operation is described in full but as the instrumented setting that makes the pharmacology measurable.

The endpoint table reorders accordingly: dose-limiting toxicity and the 3+3 escalation first, then sampling timed around the resection, then pathologic response, ctDNA clearance, and 24-month overall survival. The operation matters for one reason no metastatic study can supply: it is the only setting yielding paired pre-treatment and on-treatment human tumour tissue under controlled timing.

Measurement	Timing relative to resection	Why no other setting supplies it
Plasma pharmacokinetics	Days -14, -1, and +1	Exposure in a surgical population has not been characterised by any sponsor
Tumour pathway pharmacodynamics	Resection specimen, day 0	Cannot be obtained in the metastatic setting at all
ctDNA clearance kinetics	Days -14, +30, +90	Defines a shared surrogate anchored to a known resection date
Pathologic response grade	Resection specimen, day 0	Anchors every later perioperative RAS study to one scale

Table 7. The four measurements the consortium in application 06 exists to make, when each is taken relative to the resection, and the reason a metastatic trial cannot supply any of them.

4.2 Applications 06 to 10

06, Foundation for the NIH. A pre-competitive question no party owns: what happens pharmacologically when a RAS(ON) inhibitor is given around a curative-intent resection. The one shared output is owned by none of the three contributors, which is what makes it pre-competitive.

07, HHMI Investigator Program. Tests the one place the report's person-based argument might not reach, an investigator with no academic appointment. The programme is written as a state machine rather than a set of aims, and every failure transition re-enters it.

08, NCI Cancer Therapy Evaluation Program. The most clinically technical, because a concept review is a technical review. An adverse event here may be attributable to the drug, the operation, or the advisory system, and dose decision depends on which, so attribution is adjudicated blind before toxicity window closes.

09, Convergent Research. The only application whose organization states its own end date: no year six is budgeted, only one of three releases depends on the trial succeeding, and every artifact has a recipient class fixed in year one.

10, UC San Diego Moores Cancer Center. Different in kind: addressed to the partner institution, asking for a 45-minute feasibility meeting rather than money, and carrying three positioning corrections from the author's site research. Update: 7/29/26 30 minute meeting with UC San Diego pancreatic cancer surgeon.

4.3 What Happens If Any One Lands

The join condition is worth stating plainly. Three of the four workstreams complete without reference to which funder paid. Only the site agreement gates what follows, and no funder can supply it.




Figure 7. One award, four parallel workstreams, one join. The join condition is the site agreement, not the money, which is why nine of the ten asks can be satisfied by any single award rather than by all ten.

4.4 The Two Sets Describe One Operation

Set A and Set B are not two programmes. They describe the same hybrid procedure, carrying surgical and medical oncology arms together.

What differs is what each leads with, and that is a claim about reviewers rather than about science. A surgical reviewer asked to evaluate a drug schedule will defer to the medical oncology literature, and the reverse holds too. Writing the same trial twice, each time leading with what the reader can actually assess, is the only honest way to ask both.




Figure 8. Fourteen months of author work against the one external readout. Bar length is elapsed time, not effort, and the two 2025 overlaps are what a single operator could not have run sequentially.

5 Trial Evidence

5.1 The Chronology

In June 2025 deep-research meta-analysis over forty pancreatic cancer literature areas, more than 400,000 words, produced a ranked shortlist from which daraxonrasib emerged as the leading candidate [14]. In August 2025 a quantitative systems pharmacology simulation, ten arms and 250 ordinary differential equations per patient, returned median overall survival of 12.8 months against 5.4 for chemotherapy [4]. In October 2025 a 1000-patient digital twin returned a best progression-free survival hazard ratio of 0.31 with an ASME V&V 40 and ICH M15 aligned credibility score of 81.9 across 55 verification tests [10, 15, 16]. In May 2026, nine months after the QSP run, RASolute 302 reported 13.2 months against 6.6 in the RAS G12 population of previously treated metastatic pancreatic cancer [5]. Between June and July 2026 the regulatory package was written and two NIH funding applications were deposited within five days [17, 18].




Figure 9. Four evidence tiers drawn as containment rather than as a ladder. Nothing inside a container may be cited as though it were the container, which is the whole of the paper's evidentiary discipline.

5.2 What the Numbers Are, and What They Are Not

Source	Date	Test arm	Control	Stated limitation, in the authors' own words
Empirical triplicate, 100,000 patients	Jul 2025	G3+ AE 8.0%	G3+ AE 25.0%	KRAS G12C log file favours experimental while the KRAS-mutant report favours control; trial-to-trial variability
QSP simulation, 10 arms	Aug 2025	12.8 mo mOS, HR 0.25	5.4 mo mOS	Assumes no acquired resistance and ideal pharmacodynamics; G3+ AE consistently high in all arms
Digital twin, 1000 patients	Oct 2025	12.1 mo mOS, mPFS HR 0.31	not applicable	No patient-specific PK, PD, or tumour growth parameters; simple Emax model, no immune compartments
RASolute 302	May 2026	13.2 mo mOS	6.6 mo mOS	Metastatic and previously treated; silent on the resectable setting

Table 8. Four sources with the limitation each set of authors stated, carried in the same row as the result. No row can be read without the caveat that qualifies it, which is the condition on citing any of them.

The 2025 simulated ratio of 2.4-fold and the 2026 observed ratio of 2.0-fold are close. That proximity is a chronology observation and a hypothesis-supporting one. It is **not** a validation claim, and three differences are material: 1000 simulated against 241 enrolled; a combination against a single agent; and KRAS G12C against a primarily G12D and G12V population. Those clauses appear in all ten applications, in the same paragraph as the numbers, for the same reason they appear here.




Figure 10. Fourteen deposited works and the citation edges between them. The bill line and the trial line run in parallel and meet only twice, which is why the legislative drafts are not on the trial's path.

5.3 Cost and Schedule

Work product	Industry cost	Industry schedule	This author
Empirical virtual trial, triplicate	above \$120,000 per run	4.5 months	\$36,330 estimated, ~ one month
Ten-arm QSP virtual trial	above \$2,000,000	several months to over a year	\$36,330 estimated, ~ one month
Digital-twin simulator	\$28,000-\$700,000	3 to 16 weeks	\$36,330 estimated, ~ one month

Table 9. Three work products against their industry cost and schedule benchmarks. The author figures are estimates from the source set and are reported as such, not as audited accounts.

5.4 Provenance

Two edges cross the corridor: the verification frame from the five bill versions into the Phase 1 protocol, and the clinician objections into the principal investigator adoption guide. All else in legislation stays there.

That separation answers a question a reviewer will ask. The trial does not depend on legislation passing. The bill work supplied a verification vocabulary and a catalogue of objections, and nothing in the protocol, the IND application, or the ten applications is contingent on H. R. 9510 being enacted [19, 20].




Figure 11. Nine components, one trust boundary, and one direction of travel. Records leave for the public archive; nothing returns, and no path from the model reaches a motion controller at any point.

6 Physical AI Governance

6.1 Three Boundaries

The model has no path to the arms. Advisory output is rendered to a display. No generated token reaches a motion controller, and no software in the advisory path holds write access to the command channel. The operating surgeon approves every motion [2].

Stopping is bounded and measured. Arm-level stop is specified at 3 milliseconds and system-wide stop at 500 milliseconds, both measured on the bench before the first participant is screened and re-measured at every cohort boundary. A specification that is asserted rather than measured is not a boundary; it is a hope.

Every artifact is verified before use. The model runs on premises, pinned to a repository commit, and writes a timestamped record for every recommendation. Conversion ordered for patient safety is recorded as clinically appropriate; conversion caused by device, advisory, or integration failure counts separately against the feasibility endpoint [3].

6.2 The Operative Day

Two figures are needed for one operative day, and neither implies the other. The sequence gives the conversation, which is what an operating team recognises; the state machine gives the guards, which is what a reviewer checks. A sequence with no guards cannot say what is forbidden, and a state machine with no messages cannot say who speaks.

6.3 How a Failure Propagates

6.4 What Is Released




Figure 12. One operative day as a message sequence. Every arrow reaching the robot leaves a human first, and the two audit writes are the only messages that cross the hospital trust boundary.




Figure 13. Six advisory states and the guard on every transition. Two guards can only be satisfied by a human action, and they are the two that separate a rendered recommendation from an executed motion.




Figure 14. One top event and every path to it. Three of the four branches are AND gates, so no single failure reaches the patient, and the one OR branch is the one governed by procedure rather than by hardware.




Figure 15. Six tables and the five keys that join them. Every released record can be traced to one participant and one operative step, which is what makes an independent re-analysis of the cohort possible at all.




Figure 16. Three vendors, nine roles, and the single review chain that connects them. No model reviews its own output, and the chain is drawn as a cycle because the second reviewer’s finding returns to the author.

6.5 Who Produced Which Artifact

The documents this paper summarises were produced by a small set of frontier models with separated roles, and the separation is what makes the review chain meaningful: no model reviews its own output [6].

Model	Role	Artifacts produced
Claude Code	Author	Investigational new drug applications; trial protocols; FDA, ICH and NIH regulation adaptions; legislation revisions; machine-readable diagrams; repository platforms
ChatGPT Codex	First reviewer	Verification, validation and uncertainty quantification of author outputs; meta-analyses of trial papers; external standards; PDF reading and writing
Google Gemini	Second reviewer	Second peer review; accelerated code problem solving; meta-verification of prior stages; administrative automation

Table 10. The three frontier-model roles, taken verbatim from the repository’s own record. The separation is the point: the author model does not review itself, and the second reviewer checks the first.

7 UC San Diego Moores Cancer Center

7.1 The Positioning Constraint

UC San Diego Moores Cancer Center is named in all ten applications as the intended partner of choice, at the feasibility stage only. It is described in none of them, and not in this paper, as a partner, sponsor, trial site, or endorser, and will not be without written authorization. No agreement of any kind exists. The one application addressed to the institution asks for a 45-minute meeting.

The constraint is not a disclaimer bolted on at the end; it shapes what the applications can claim. A funder reading application 01 is told the site is sought, not secured, and the first-year gate reads *site executed* rather than *site confirmed*. An application that overstated the relationship would fail its own go or no-go condition in the first year.

This section carries no figure. Application 10 already contains the intake sequence, the authority diagram, the feasibility grid, and the approval dependency, and drawing any of them again would be duplication rather than a new perspective. The requirement that every figure be a unique perspective is a constraint on adding figures, not a quota to fill.

7.2 Three Corrections the Site Research Forced

The most useful thing the site research produced was not a contact list. It was a list of things the applications were getting wrong.

Claim as first drafted	Correction	Where it now appears
Daraxonrasib is first-in-human	It is investigational and already in Phase 3 evaluation. The supportable claim is first prospective clinical evaluation of the integrated surgical and advisory workflow, subject to FDA & institutional confirmation	A positioning note in nine applications; a pre-send checklist line in all ten
An eight-arm robotic Whipple	No site can assess feasibility or regulatory status without manufacturer, model, cleared configuration, arm count, instruments, and software version. The configuration is specified at the site agreement	Application 10 §1; the positioning note in the other nine
Drug access is available	No drug supply, letter of authorization, or regulatory cross-reference is in place with the agent's developer	Application 06; application 10 §1

Table 11. Three corrections the site research forced, and where each now appears. Nine of the ten applications carried the first error until the PART I audit pass, which is recorded here rather than quietly fixed.

The first correction is the kind of error that survives review. *First-in-human* was true of the integrated workflow and false of the drug, and in a title the two readings are indistinguishable. A reviewer at a drug-development office would have caught it at once and discounted everything after it. The correction cost two sentences; not making it would have cost the application.

7.3 Two Routes, Sequenced Differently

The repository carries research on two candidate sites, and are sequenced in opposite directions on purpose.

	UC San Diego	Scripps
Entry point	Pancreatic surgery, GI med. oncology	AI validation and digital-trial methods
First ask	A feasibility meeting on the trial	Independent review of AI portfolio
The trial appears at	Step 3 of 12	Step 9 of 13
Risk if the route fails	The trial has no site	AI work has no independent reviewer

Table 12. The two site routes on four properties. They are sequenced in opposite directions because they are asking different institutions for different things, not because one is a fallback for the other.

The UC San Diego route is trial-first because Moores has the pancreatic surgical volume and trial operations the study needs. The Scripps route is AI-first because the value it can add is independent evaluation of verification methods, and asking a methods group to host a surgical trial would waste the introduction.

Application 03 uses the contrast directly: its pathway figure shows the branch taken if the first site agreement does not execute, and that branch is a second qualified site rather than the end of the programme, because the performer, not the site, holds the regulatory package.

7.4 What a Productive Early Outcome Looks Like

The site research sets a deliberately modest bar and the applications inherit it. A productive outcome is an invitation to present, a clinical principal investigator identified, acceptance into a feasibility process, a preliminary budget and accrual estimate, a technical review, a defined regulatory pathway, or support for a grant application. None of those is a trial, and none is described as one.

A patient-facing trial proceeds only after the site, the manufacturer, the FDA pathway, the funding, the sponsor responsibilities, the contracts, the insurance, the review board approval, and the operational resources are all documented. That sentence is in the site research, in every application's back matter, and here, and it is the same sentence in all three places.

8 Budget and Leverage

8.1 One Budget, Ten Shapes

The underlying cost is one number: \$700,000 per year of direct cost, \$3,500,000 over five years, no cost share, no programme income anticipated [18]. What changes across the ten is the shape that number takes, because a mechanism is in part a claim about time.




Figure 17. Five years of cash on the left, contributed non-federal value on the right, and the two layers that appear on both sides. The contributed column is what the annex asks agencies to prioritise.

No	Recipient	Term	Direct ask	Shape, and why the mechanism requires it
01	NIH Pioneer Award	5 years	\$3,500,000	Fully funded in year one, as the annex asks for long-duration grants
02	ARPA-H	36 months	\$2,100,000	Compressed to three gates, because an ARPA bet should reach a defensible dose or stop
03	NSF TIP X-Labs	5 years	\$3,500,000	One public good released per year, so the organization is judged annually
04	DOE Genesis Mission	5 years	\$3,500,000	Mission artifacts released in years 1, 2 to 3, and 3 to 5
05	NIH SEED SBIR	9 + 24 months	\$1,606,000	Split at a measured gate: \$306,000 Phase I, \$1,300,000 Phase II
06	FNIH AMP	5 years	\$3,500,000 cash	Cash only; drug, theatre, pathology and assay support are contributed
07	HHMI Investigator	7 years	\$700,000 per year	Clinical cost only; the balance funds the post-readout pharmacology
08	NCI CTEP	5 years	\$3,500,000	Correlative work costed inside, not deferred to a supplement
09	Convergent FRO	5 years	\$3,500,000	No year six is budgeted, and year five carries closure costs explicitly
10	UCSD Moores	one meeting	none	The only ask that is not money

Table 13. One underlying budget in ten shapes. The mechanism determines the shape: a gated agency compresses the term, a small-business route splits it at a measured gate, and a time-bound organization budgets its own closure.




Figure 18. Ten submissions sent on one day and ten review clocks that are not the same length. The site milestones beneath them are the binding constraint: no award can start before the site agreement executes.

8.2 Leverage

The annexed FY 2028 memorandum asks agencies to “expand the use of non-Federal cost share, so that Federal dollars draw in and are amplified by private and other non-Federal investment rather than standing alone” [11]. Application 06 is written to that instruction: the cash ask covers clinical conduct and regulatory work, and drug supply, theatre time, pathology, and bioanalytical support are contributed.

The contributed column carries no dollar figure, deliberately. No agreement exists with any developer or site, so a valuation would be invented, and an invented cost-share figure is worse than none: it converts a real structural argument into an unverifiable number. The claim made is qualitative and checkable, that four named categories of cost would not appear in the federal request.

8.3 The Binding Constraint Is Not Money

Seven of the ten review clocks close before the projected site-agreement date and three run past it. The programme’s schedule is therefore governed by the institution rather than by any funder: an award arriving in month four cannot begin work until month six at the earliest, and the two regulatory activities identified in §2 cannot start until the site is executed however much money is in hand.

That is why application 10, the one with no financial ask, is the most consequential of the ten. It is the only one addressed to the party that controls the constraint.

9 Build Method

9.1 Two Schedules, Thirteen Sub-Prompts

The whole build runs from one master prompt, which generates thirteen stage sub-prompts in two schedules that never share a stage: five drive the ten application file sets, eight drive this paper. Each names its output directory, its commit floor, and the repository files it may read.

Part	No	Stage	Output	Figures
PART I	1	Recipients and templates	Recipient list, cover variants, style and bib contract	none
PART I	2	Set A, surgical	Applications 01 to 05	17
PART I	3	Set B, medical oncology	Applications 06 to 10	17
PART I	4	Email, attached	Ten ‘.txt’ files, ten Overleaf zips	none
PART I	5	READMEs, audit	Every ‘funding/**’ README; the error-fix pass	none
PART II	1	mermaid	Six figure specifications	6
PART II	2	plantuml	Three figure specifications	3
PART II	3	d2	Four figure specifications	4
PART II	4	diagrams, python	Three figure specifications	3
PART II	5	graphviz	Four figure specifications	4
PART II	6	draft-apply	Skeleton, drafting instructions, 20 figure slots	20 placed
PART II	7	full-apply	Every instruction resolved, every figure drawn	20 drawn
PART II	8	final-apply	Senior-author proof-reading pass	20 polished

Table 14. The thirteen stage sub-prompts in two schedules. The application stages carry 34 figures across ten attachments; the paper stages carry 20, specified once and drawn twice.

9.2 Why the Diagram Split Is Uneven

An equal split across five platforms would be four figures each. The split used here is 6, 3, 4, 3, 4, and the difference is not arbitrary.

Platform	Count	The question only this vocabulary can state
Mermaid-type	6	What happens next, what decides, how long, and who spoke in what order
PlantUML-type	3	Who can act, under what guard a state changes, and what runs concurrently
D2-type	4	What contains what, what tabulates against what, and what joins to what
Diagrams (python)-type	3	What runs where, across which trust boundary, and on whose hardware
Graphviz-type	4	What depends on what, how a failure propagates, what one record answers

Table 15. The five platforms, their share of the twenty figures, and the question each is the only one able to state. Forcing an equal split would have moved at least four figures into a vocabulary that cannot carry their subject.

Mermaid takes the largest share because the paper’s spine is chronological: a policy document, ten submissions, a fourteen-month record, and one operative day are all sequences. PlantUML takes the smallest because only three subjects need a guard grammar. Forced to four each, the authority diagram would have become a flowchart, which can’t draw the absence of a permission; the evidence tiers a ladder, which asserts order where there is containment; and fault tree a dependency graph, and can’t discern AND from OR.




Figure 19. Five layers, executed twice. The left column is one operator with an LLM toolchain and a public repository; the right is the institutional equivalent. Only the bottom layer is common to both.

9.3 One Operator Against a Program Office

9.4 What Is Reproducible, and What Is Not

Reproducible: the master prompt, the thirteen sub-prompts, the named sources, the commit sequence, the figure construction notes, and the compile command. A reader with the repository can rebuild every artifact and check every number against the file it came from.

Not reproducible: the judgement at each step. Which recipient qualifies, which vocabulary a figure needs, which sentence in a 123-page report is worth quoting, and which of the author's own claims is wrong, these are decisions, and a second run would make some differently. That is why the author's review is a required stage rather than an optional one, and why the PART I audit pass, which found the first-in-human error in nine of ten applications, is recorded in §7 rather than absorbed silently.




Figure 20. Five verification artifacts and the reviewer question each field answers. Every field is released with the cohort, and the two marked fields are the ones that can falsify the programme's central claim.

10 Risks and Limits

10.1 What Is Not Claimed

This section is written against the paper rather than for it.

Each row pairs an inference a reader might reasonably draw with what is actually true.

Inference a reader might draw	What is actually true
The simulations validate RASolute 302	They preceded it. The proximity is a chronology observation, with differences: KRAS G12C against a primarily G12D and G12V population
Daraxonrasib is first-in-human	It is investigational and already in Phase 3 evaluation. The supportable claim concerns the integrated surgical and advisory workflow
A robotic configuration has been specified or cleared	None has. Manufacturer, model, cleared configuration, arm count, instruments and software version are named at the site agreement
Drug supply or a cross-reference exists	Neither exists. No agreement of any kind is in place with the agent's developer
UC San Diego has agreed to something	It has not. One application asks for a 45-minute feasibility meeting
A patient has been treated	None has. The trial does not exist until a site, an IRB, an active IND, monitoring, insurance, and a sponsor of record are all in place
The advisory system holds a regulatory clearance	It holds none, and the applications ask for the study that would generate the evidence a clearance would require

Table 16. Seven inferences a reader might reasonably draw, and what is actually true in each case. Every row corresponds to a sentence carried in the back matter of all ten applications.

10.2 What Would Falsify the Case

A claim that cannot fail is not a claim. Two released fields can falsify the programme's central assertion, and both are released whatever they show.

The first is `credibility_score`. The digital-twin work carries a score of 81.9 under ASME V&V 40 and ICH M15 alignment across 55 verification notebook tests [10, 15, 16]. If an independent run does not reproduce it, the modelling tier loses its standing and with it the arm-selection rationale.

The second is `negative_results`. The applications commit to releasing negative and inconclusive cohorts, and §6's Gold Standard Science mapping makes that a field rather than a promise [21]. A released cohort with that field empty falsifies claim to Gold Standard alignment directly, whatever clinical result.

10.3 Risks the Programme Cannot Retire Alone

Risk	Who can retire it	What the applicant can do meanwhile
No qualified site	An academic cancer center	Keep regulatory package with performer, so a second site substitutes without a protocol change
No drug supply or cross-reference	The agent's developer	State plainly in every document that none exists, and design request so it can be answered once
No sponsor-investigator determination	Institutional contracting and regulatory offices	Ask for the determination before the IRB submission, not after
Robotic configuration not permissible	The site's surgical robotics programme	Leave the configuration unspecified until the site names it
Accrual below plan	The site's referral base	Extend accrual window vs. widening eligibility

Table 17. Five risks the programme cannot retire on its own, who can retire each, and the action available meanwhile. None of the five is retired by funding, which is why the funding argument and the feasibility argument are separate.

10.4 The Limit of the Whole Exercise

Ten applications on one piece of science demonstrate that the mechanism is a variable an individual applicant can address. They do not demonstrate that the science is right, that the operation is safe, or that the drug helps in the resectable setting. Those questions are settled by the trial, and the trial does not exist. What this paper reports is the state of the argument on August 4, 2026, with every number traceable to a named file and every claim paired with the limitation its authors stated.

Back Matter

Abbreviations

Short	Expansion	Short	Expansion
AE	Adverse event	MTD	Maximum tolerated dose
AMP	Accelerating Medicines Partnership	NCI	National Cancer Institute
ARPA-H	Advanced Research Projects Agency for Health	NIH	National Institutes of Health
CTCAE	Common Terminology Criteria for Adverse Events	NSF	National Science Foundation
CTEP	Cancer Therapy Evaluation Program	ODE	Ordinary differential equation
ctDNA	Circulating tumour DNA	OS	Overall survival
DLT	Dose-limiting toxicity	PD	Pharmacodynamics
DOI	Digital object identifier	PDAC	Pancreatic ductal adenocarcinoma
FNIH	Foundation for the National Institutes of Health	PFS	Progression-free survival
FRO	Focused research organization	PK	Pharmacokinetics
HHMI	Howard Hughes Medical Institute	QSP	Quantitative systems pharmacology
IIT	Investigator-initiated trial	RP2D	Recommended Phase 2 dose
IND	Investigational new drug	SAE	Serious adverse event
IRB	Institutional review board	SBIR	Small Business Innovation Research
ISGPS	International Study Group of Pancreatic Surgery	TIP	Technology, Innovation and Partnerships
LLM	Large language model	VVUQ	Verification, validation and uncertainty quantification

Table 18. Abbreviations used in this paper, set in two column pairs at the body measure. Trial-specific terms follow the definitions in the published Phase 1 protocol rather than being redefined here.

Data and Code Availability

Every work cited to a digital object identifier is openly deposited. The repository is at <https://github.com/kevinkawchak/physical-ai-oncology-trials>, version v4.4.0 [6].

It contains the ten application file sets with their email text and Overleaf bundles, the twenty figure specifications across five machine-readable diagram platforms, the three build stages of this paper, the master prompt and the thirteen sub-prompts, and the TikZ source of every figure printed here. No raster image is used anywhere in this paper or in any of the ten attachments.

Author Contributions and Conflicts

Kevin Kawchak is the sole author. He is chief executive of ChemicalQDevice and holds no equity in Revolution Medicines or in any robotic surgery vendor. No sponsor, contract research organization, institutional review board, regulator, or medical society reviewed this paper or authorised any statement in it. The work was adapted using Claude Code Opus 5.

Citation

Kawchak K. 10 Funding Applications: A Phase 1, First-in Human, PDAC Clinical Trial Protocol of a LLM-Directed Robotic Whipple with Daraxonrasib (RMC-6236). Zenodo. 2026; doi: [10.5281/zenodo.21787424](https://doi.org/10.5281/zenodo.21787424).

References

- [1] Michael Kratsios. Science: A New Golden Age. Report to the President of the United States, July 2026. URL <https://www.whitehouse.gov/science/>. Letter of Transmittal, July 21, 2026.
- [2] Kevin Kawchak. On-Premises LLM-Directed Robotic Pancreaticoduodenectomy with Perioperative Daraxonrasib (RMC-6236) in KRAS-Mutated Pancreatic Ductal Adenocarcinoma. Zenodo, June 2026. URL <https://doi.org/10.5281/zenodo.20780121>.
- [3] Kevin Kawchak. Investigational New Drug Application: Daraxonrasib, Phase 1, LLM-Directed Robotic Whipple in KRAS-Mutated PDAC. Zenodo, July 2026. URL <https://doi.org/10.5281/zenodo.21097442>.
- [4] Kevin Kawchak. QSP Metastatic Pancreatic Cancer AI Clinical Trial Simulation: From Protocol to Prediction, Code, VVUQ, and Playbook. Zenodo, August 2025. URL <https://doi.org/10.5281/zenodo.17001137>.
- [5] Eileen M. O'Reilly et al. Daraxonrasib or Chemotherapy in Previously Treated Metastatic Pancreatic Cancer. *New England Journal of Medicine*, May 2026. doi: 10.1056/NEJMoa2605555. URL <https://doi.org/10.1056/NEJMoa2605555>.
- [6] Kevin Kawchak. physical-ai-oncology-trials, repository v4.4.0. GitHub, 2026. URL <https://github.com/kevinkawchak/physical-ai-oncology-trials>.
- [7] Jack W. Scannell, Alex Blanckley, Helen Boldon, and Brian Warrington. Diagnosing the Decline in Pharmaceutical R&D Efficiency. *Nature Reviews Drug Discovery*, 11:191–200, 2012. doi: 10.1038/nrd3681. URL <https://doi.org/10.1038/nrd3681>. Eroom's law, cited by the report for the productivity slowdown.
- [8] Pierre Azoulay, Joshua S. Graff Zivin, and Gustavo Manso. Incentives and Creativity: Evidence from the Academic Life Sciences. *The RAND Journal of Economics*, 42(3):527–554, 2011. doi: 10.1111/j.1756-2171.2011.00140.x. URL <https://doi.org/10.1111/j.1756-2171.2011.00140.x>. The HHMI-style person-based funding comparison cited by the report.
- [9] Kevin Kawchak. Physical AI Oncology Trial Founding Documents. ChemicalQDevice repository, July 2026. URL <https://github.com/kevinkawchak/physical-ai-oncology-trials/blob/main/funding/supplementary/Physical%20AI%20Oncology%20Trial%20Founding%20Documents.md>. Fourteen prior works across funding, trial, and legislation.
- [10] Kevin Kawchak. Daraxonrasib Efficient LLM Trial Simulations. ChemicalQDevice repository source set, July 2026. URL <https://github.com/kevinkawchak/physical-ai-oncology-trials/blob/claude/phase-1-pancreatic-funding-apps-t3jcad/funding/supplementary/source-files/Daraxonrasib-Efficient-LLM-Trial-Simulations.zip>.
- [11] Michael Kratsios and Russell T. Vought. Fiscal Year 2028 Administration Research and Development Budget Priorities. Joint OSTP and OMB memorandum NSTM-5 / M-26-16, July 2026. URL <https://www.whitehouse.gov/wp-content/uploads/2026/07/White-House-Fiscal-Year-2028-RD-Priorities-Memorandum.pdf>. Annex to Science: A New Golden Age; names the six national missions.
- [12] Claudio Bassi, Giovanni Marchegiani, Christos Dervenis, et al. The 2016 Update of the International Study Group (ISGPS) Definition and Grading of Postoperative Pancreatic Fistula. *Surgery*, 161(3):584–591, 2017. doi: 10.1016/j.surg.2016.11.014. URL <https://doi.org/10.1016/j.surg.2016.11.014>.
- [13] Executive Order No. 14363, Launching the Genesis Mission, November 2025. URL <https://www.federalregister.gov/document/s/2025/11/28/2025-21665/launching-the-genesis-mission>.
- [14] Kevin Kawchak. End-to-End Oncology Clinical Trial LLM Efficiency: Daraxonrasib Identification. Zenodo, June 2025. URL <https://doi.org/10.5281/zenodo.15735068>.
- [15] The American Society of Mechanical Engineers. Verification, Validation and Uncertainty Quantification (VVUQ). ASME standard, 2026. URL <https://www.asme.org/codes-standards/publications-information/verification-validation-uncertainty>.
- [16] U.S. Food and Drug Administration. M15 General Principles for Model-Informed Drug Development. GUIDANCE DOCUMENT, 2026. URL <https://www.fda.gov/regulatory-information/search-fda-guidance-documents/m15-general-principles-model-informed-drug-development>.
- [17] Kevin Kawchak. Clinical Trial Funding Application, RFA-RM-27-001. Zenodo, July 2026. URL <https://doi.org/10.5281/zenodo.21232965>.
- [18] Kevin Kawchak. Clinical Trial Funding Application v2.0, RFA-RM-27-001. Zenodo, July 2026. URL <https://doi.org/10.5281/zenodo.21317266>.
- [19] Kevin Kawchak. H. R. 9510 (Bill v5.0) 2026. Zenodo, June 2026. URL <https://doi.org/10.5281/zenodo.20619762>.
- [20] Kevin Kawchak. Earning the Clinician's Trust: New Framework for Verified Physical AI Oncology Trials. Zenodo, June 2026. URL <https://doi.org/10.5281/zenodo.20710602>.
- [21] Executive Order No. 14303, Restoring Gold Standard Science, May 2025. URL <https://www.federalregister.gov/documents/2025/05/29/2025-09802/restoring-gold-standard-science>. Nine principles binding on all federally funded research.